

p8106 Final Project

Qixiang Chen

5/5/2022

TO-DO List: 1. explanatory analysis (Is there any interesting structure present in the data? What were your findings?) 2. visualization work (feature plot done, plots for factors) 3. build 6 classification models 4. build a frame for the final model selection (ROC, AUC, ConfusionMatrix, error rate) 5. Try to list out the important predictors 6. tuning parameters 7. interpretation of each model

In this dataset, **age** refers to age in days. For variable **gender**, 1 represents women, 2 represents men. Thus, we need to do the corresponding adjustments to make it look formal.

1. Since variable **id** does not contribute to the following analysis, we exclude **id** from the dataset.
2. For variable **gender**, 1 represents women, 2 represents men. To make it serve as a dummy variable, we convert it into factor.

```
df = read.csv("./cardio_train.csv", header = TRUE, stringsAsFactors = FALSE, sep = ";") %>%
  janitor::clean_names() %>%
  dplyr::select(-id) %>%
  rename(age_day = age) %>%
  mutate(gender = gender - 1,
         cholesterol = case_when(cholesterol == 1 ~ "normal",
                                cholesterol == 2 ~ "above normal",
                                cholesterol == 3 ~ "well above normal"
                                ),
         gluc = case_when(gluc == 1 ~ "normal",
                          gluc == 2 ~ "above normal",
                          gluc == 3 ~ "well above normal"
                          ),
         gender = as.factor(gender),
         smoke = as.factor(smoke),
         alco = as.factor(alco),
         active = as.factor(active),
         cardio = case_when(cardio == 0 ~ "nondiseased",
                             cardio == 1 ~ "diseased"
                             ),
         cardio = as.factor(cardio)
  ) %>%
  mutate(cholesterol = factor(cholesterol, levels = c("normal", "above normal", "well above normal")),
         gluc = factor(gluc, levels = c("normal", "above normal", "well above normal")),
         cardio = factor(cardio, levels = c("nondiseased", "diseased"))
  ) %>%
  dplyr::select(age_day, height, weight, ap_hi, ap_lo, everything())

#df
```

```
#Check whether there is any missing value.
```

```
missing_train = sapply(df, function(x) sum(is.na(x)))  
print(missing_train[missing_train > 0])
```

```
## named integer(0)
```


```
#Summary
```

```
summary(df)
```

```
##      age_day      height      weight      ap_hi  
##  Min.   :10798  Min.   : 55.0  Min.   : 10.00  Min.   : -150.0  
## 1st Qu.:17664 1st Qu.:159.0 1st Qu.: 65.00 1st Qu.: 120.0  
## Median :19703 Median :165.0 Median : 72.00 Median : 120.0  
## Mean   :19469 Mean   :164.4 Mean   : 74.21 Mean   : 128.8  
## 3rd Qu.:21327 3rd Qu.:170.0 3rd Qu.: 82.00 3rd Qu.: 140.0  
## Max.   :23713 Max.   :250.0 Max.   :200.00 Max.   :16020.0  
##      ap_lo      gender      cholesterol  
##  Min.   : -70.00  0:45530  normal      :52385  
## 1st Qu.:  80.00  1:24470  above normal : 9549  
## Median :  80.00      well above normal: 8066  
## Mean   :  96.63  
## 3rd Qu.:  90.00  
## Max.   :11000.00  
##      gluc      smoke      alco      active      cardio  
## normal      :59479  0:63831  0:66236  0:13739  nondiseased:35021  
## above normal : 5190  1: 6169  1: 3764  1:56261  diseased   :34979  
## well above normal: 5331  
##  
##  
##
```

```
#Corr Plot
```

```
x_df = model.matrix(cardio ~ ., df)[-1]  
corrplot::corrplot(cor(x_df),  
                    method = "circle",  
                    type = "full",  
                    tl.cex = 0.5)
```



```

set.seed(2022)
# Randomly sample 3500 data points without replacement from the data set.
df_sample = sample_n(df, 1000) %>%
  janitor::clean_names()

#colnames(df_sample)[0] <- "id"

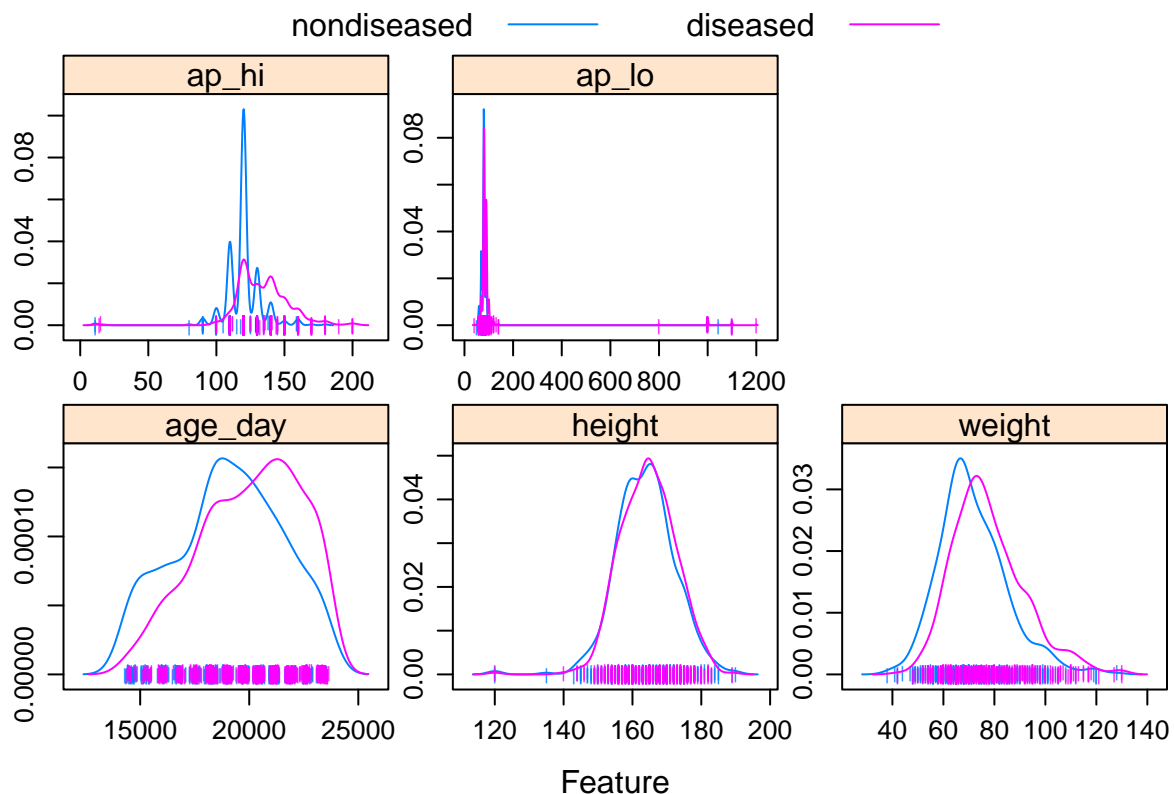
#save R data
save(df_sample, file = "df_sample.RData")

#summary(df_sample)

load("df_sample.RData")

#Feature Plot
featurePlot(x = df_sample[, 1:5],
            y = df_sample$cardio,
            scales = list(x = list(relation="free"),
                           y = list(relation="free")),
            plot = "density",
            pch = "|",
            auto.key = list(columns = 2)
            )

```



```
set.seed(2022)
training_tag = createDataPartition(y = df_sample$cardio,
                                   p = 0.7,
                                   list = FALSE)

# For training dataset
training_data = df_sample[training_tag, ] %>% janitor::clean_names()
training_predictors_x = model.matrix(cardio ~ ., training_data)[, -1]
training_outcome_y = training_data$cardio

# For test dataset
test_data = df_sample[-training_tag, ] %>% janitor::clean_names()
test_predictors_x = model.matrix(cardio ~ ., test_data)[, -1]
test_outcome_y = test_data$cardio

# Control
control = trainControl(method = "repeatedcv",
                       summaryFunction = twoClassSummary,
                       repeats = 5,
                       classProbs = TRUE)
```

1.logistic regression

```
#undiseased: 0
#diseased: 1
contrasts(df_sample$cardio)
```

```
##           diseased
## nondiseased      0
## diseased         1
```

```
set.seed(2022)
glm_fit = glm(cardio ~ .,
              data = df_sample,
              subset = training_tag,
              family = binomial(link = "logit")
            )
summary(glm_fit)
```

```
##
## Call:
## glm(formula = cardio ~ ., family = binomial(link = "logit"),
##      data = df_sample, subset = training_tag)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5713  -0.9408   0.0311   0.9972   3.2204
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -9.663e+00  2.236e+00  -4.322 1.54e-05 ***
## age_day         1.501e-04  3.707e-05   4.049 5.13e-05 ***
## height        -5.518e-03  1.213e-02  -0.455  0.64931
## weight         1.377e-02  6.885e-03   2.000  0.04555 *
## ap_hi          5.216e-02  6.995e-03   7.457 8.86e-14 ***
## ap_lo          1.214e-03  9.387e-04   1.293  0.19606
## gender1        1.335e-01  2.127e-01   0.627  0.53036
## cholesterolabove normal  1.997e-01  2.708e-01   0.738  0.46081
## cholesterolwell above normal 8.861e-01  3.318e-01   2.671 0.00757 **
## glucabove normal -2.142e-02  3.362e-01  -0.064  0.94919
## glucwell above normal  9.407e-02  3.552e-01   0.265  0.79112
## smoke1        -4.787e-01  3.356e-01  -1.426  0.15376
## alco1         -2.794e-01  4.308e-01  -0.648  0.51670
## active1       -2.280e-01  2.058e-01  -1.108  0.26785
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 970.41  on 699  degrees of freedom
## Residual deviance: 804.40  on 686  degrees of freedom
## AIC: 832.4
##
## Number of Fisher Scoring iterations: 5
```

```

test_pred_prob = predict(glm_fit, newdata = test_data,
                          type = "response"
                        )

test_pred = rep("nondiseased", length(test_pred_prob))

test_pred[test_pred_prob > 0.5] = "diseased"

confusionMatrix(data = as.factor(test_pred),
                 reference = test_outcome_y
                )

```

```

## Warning in confusionMatrix.default(data = as.factor(test_pred), reference
## = test_outcome_y): Levels are not in the same order for reference and data.
## Refactoring data to match.

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  nondiseased diseased
##  nondiseased      111      34
##   diseased       39      116
##
##              Accuracy : 0.7567
##              95% CI : (0.704, 0.8041)
##   No Information Rate : 0.5
##   P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.5133
##
##  Mcnemar's Test P-Value : 0.6397
##
##              Sensitivity : 0.7400
##              Specificity : 0.7733
##              Pos Pred Value : 0.7655
##              Neg Pred Value : 0.7484
##              Prevalence : 0.5000
##              Detection Rate : 0.3700
##              Detection Prevalence : 0.4833
##              Balanced Accuracy : 0.7567
##
##              'Positive' Class : nondiseased
##

```

```

auc(test_outcome_y, test_pred_prob)

```

```

## Setting levels: control = nondiseased, case = diseased

```

```

## Setting direction: controls < cases

```

```

## Area under the curve: 0.8053

```

```
#caret logistic for model selection
set.seed(2022)
logistic_caret = train(x = training_predictors_x,
                        y = training_outcome_y,
                        method = "glm",
                        metric = "ROC",
                        trControl = control
                        )
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic_caret)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5713  -0.9408   0.0311   0.9972   3.2204
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -9.663e+00  2.236e+00  -4.322 1.54e-05 ***
## age_day         1.501e-04  3.707e-05   4.049 5.13e-05 ***
## height        -5.518e-03  1.213e-02  -0.455  0.64931
## weight         1.377e-02  6.885e-03   2.000  0.04555 *
## ap_hi          5.216e-02  6.995e-03   7.457 8.86e-14 ***
## ap_lo          1.214e-03  9.387e-04   1.293  0.19606
## gender1        1.335e-01  2.127e-01   0.627  0.53036
## `cholesterolabove normal` 1.997e-01  2.708e-01   0.738  0.46081
## `cholesterolwell above normal` 8.861e-01  3.318e-01   2.671  0.00757 **
## `glucabove normal`    -2.142e-02  3.362e-01  -0.064  0.94919
## `glucwell above normal` 9.407e-02  3.552e-01   0.265  0.79112
## smoke1         -4.787e-01  3.356e-01  -1.426  0.15376
## alco1         -2.794e-01  4.308e-01  -0.648  0.51670
## active1        -2.280e-01  2.058e-01  -1.108  0.26785
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 970.41  on 699  degrees of freedom
## Residual deviance: 804.40  on 686  degrees of freedom
## AIC: 832.4
##
## Number of Fisher Scoring iterations: 5
```

2. MARS

```
#adjust cardio to be dummy
```


```
set.seed(2022)
mars_model = train(x = training_predictors_x,
                   y = training_outcome_y,
                   method = "earth",
                   tuneGrid = expand.grid(degree = 1:3, nprune = 2:13),
                   metric = "ROC",
                   trControl = control
                   )
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


[illegible]

```
## Call: earth(x=matrix[700,13], y=factor.object, keepxy=TRUE,
##           glm=list(family=function.object, maxit=100), degree=2, nprune=7)
##
## GLM coefficients
##
##              diseased
## (Intercept)      -2.85083364
## cholesterolwell above normal    2.29450622
## h(16739-age_day)      -0.00034732
## h(ap_hi-100)          0.09199229
## h(ap_hi-100) * cholesterolwell above normal -0.05256094
## h(age_day-16739) * h(ap_hi-150)      -0.00000281
## h(age_day-16739) * h(150-ap_hi)      0.00000620
##
## GLM (family binomial, link logit):
## nulldev df      dev df   devratio    AIC iters converged
##  970.406 699   778.558 693     0.198   792.6     5         1
##
## Earth selected 7 of 21 terms, and 3 of 13 predictors (nprune=7)
## Termination condition: Reached nk 27
## Importance: ap_hi, age_day, cholesterolwell above normal, height-unused, ...
## Number of terms at each degree of interaction: 1 3 3
## Earth GCV 0.1996215   RSS 133.4202   GRSq 0.2037937   RSq 0.237599
```

```
ggplot(mars_model, highlight = T)
```



```
mars_model$bestTune
```

```
##      nprune degree
## 18         7      2
```

```
mars_model$results
```


##	degree	nprune	ROC	Sens	Spec	ROCSD	SensSD	SpecSD
## 1	1	2	0.7393469	0.7800000	0.5971429	0.05727936	0.07982697	0.08909747
## 13	2	2	0.7394449	0.7800000	0.5971429	0.05720819	0.07982697	0.08909747
## 25	3	2	0.7396816	0.7800000	0.5971429	0.05739157	0.07982697	0.08909747
## 2	1	3	0.7615184	0.7554286	0.6342857	0.05905114	0.08870581	0.09537405
## 14	2	3	0.7595347	0.7285714	0.6571429	0.05778705	0.10740977	0.09361098
## 26	3	3	0.7616000	0.7222857	0.6640000	0.06089484	0.09881261	0.09788433
## 3	1	4	0.7590776	0.7702857	0.6200000	0.06109256	0.09052763	0.11504848
## 15	2	4	0.7610694	0.7177143	0.6668571	0.05688798	0.09460939	0.09389707
## 27	3	4	0.7594531	0.7194286	0.6531429	0.06204869	0.09531989	0.09641815
## 4	1	5	0.7613878	0.7720000	0.6188571	0.05886146	0.08820294	0.11379047
## 16	2	5	0.7638122	0.7240000	0.6548571	0.05750386	0.08727634	0.09499600
## 28	3	5	0.7649143	0.7251429	0.6588571	0.06261853	0.09215636	0.09904332
## 5	1	6	0.7605551	0.7760000	0.6137143	0.05871067	0.08836711	0.09799999
## 17	2	6	0.7674041	0.7285714	0.6640000	0.06010477	0.08468773	0.08649980
## 29	3	6	0.7691673	0.7285714	0.6594286	0.06146664	0.08449078	0.10483362
## 6	1	7	0.7595347	0.7702857	0.6154286	0.05501890	0.08638417	0.09782474
## 18	2	7	0.7701388	0.7302857	0.6640000	0.05934579	0.07875753	0.09317573
## 30	3	7	0.7691510	0.7251429	0.6611429	0.06378316	0.08579005	0.10114052
## 7	1	8	0.7562204	0.7708571	0.6165714	0.05571805	0.08173259	0.09982742
## 19	2	8	0.7684490	0.7314286	0.6594286	0.05427342	0.08224262	0.08976623
## 31	3	8	0.7672653	0.7274286	0.6605714	0.06017651	0.08349109	0.09566010
## 8	1	9	0.7579020	0.7720000	0.6205714	0.05795684	0.08571234	0.09998584
## 20	2	9	0.7669714	0.7348571	0.6588571	0.05782721	0.08385545	0.09042636
## 32	3	9	0.7655347	0.7251429	0.6582857	0.06142293	0.07954263	0.09846130
## 9	1	10	0.7566367	0.7725714	0.6234286	0.05599165	0.08162449	0.09858474
## 21	2	10	0.7650857	0.7331429	0.6594286	0.05806574	0.08337128	0.09482046
## 33	3	10	0.7630857	0.7234286	0.6582857	0.06324309	0.07775269	0.09997251
## 10	1	11	0.7566204	0.7708571	0.6240000	0.05497303	0.08374611	0.09699526
## 22	2	11	0.7644163	0.7337143	0.6588571	0.05822727	0.08472510	0.09474488
## 34	3	11	0.7637551	0.7251429	0.6582857	0.06189125	0.07501548	0.09863035
## 11	1	12	0.7570122	0.7708571	0.6240000	0.05451333	0.08374611	0.09699526
## 23	2	12	0.7651673	0.7354286	0.6588571	0.05762884	0.08538518	0.09474488
## 35	3	12	0.7644735	0.7274286	0.6588571	0.06146464	0.07552896	0.09751771
## 12	1	13	0.7570122	0.7708571	0.6240000	0.05451333	0.08374611	0.09699526
## 24	2	13	0.7649878	0.7337143	0.6594286	0.05778366	0.08781488	0.09517121
## 36	3	13	0.7644735	0.7274286	0.6588571	0.06146464	0.07552896	0.09751771

```
coef(mars_model$finalModel)
```

```
##              (Intercept)
##              -2.850834e+00
##              h(ap_hi-100)
##              9.199229e-02
##              h(16739-age_day)
```

```
## -3.473246e-04
## h(age_day-16739) * h(ap_hi-150)
## -2.808598e-06
## h(age_day-16739) * h(150-ap_hi)
## 6.204302e-06
## cholesterolwell above normal
## 2.294506e+00
## h(ap_hi-100) * cholesterolwell above normal
## -5.256094e-02
```

```
vip(mars_model$finalModel)
```



```
mars_test_pred_prob_df = predict(mars_model, newdata = test_predictors_x,
                                  type = "prob"
                                )

mars_test_pred_prob = mars_test_pred_prob_df$diseased

mars_test_pred = rep("nondiseased", length(mars_test_pred_prob))

mars_test_pred[mars_test_pred_prob > 0.5] = "diseased"

confusionMatrix(data = as.factor(mars_test_pred),
                 reference = test_outcome_y
                )
```

```
## Warning in confusionMatrix.default(data = as.factor(mars_test_pred), reference
## = test_outcome_y): Levels are not in the same order for reference and data.
## Refactoring data to match.
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  nondiseased diseased
## nondiseased      109      31
## diseased         41     119
##
##           Accuracy : 0.76
##           95% CI : (0.7076, 0.8072)
##       No Information Rate : 0.5
##       P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.52
##
##  Mcnemar's Test P-Value : 0.2888
##
##           Sensitivity : 0.7267
##           Specificity : 0.7933
##       Pos Pred Value : 0.7786
##       Neg Pred Value : 0.7438
##           Prevalence : 0.5000
##       Detection Rate : 0.3633
##   Detection Prevalence : 0.4667
##       Balanced Accuracy : 0.7600
##
##       'Positive' Class : nondiseased
##
```

```
auc(test_outcome_y, mars_test_pred_prob)
```

```
## Area under the curve: 0.8135
```


```
coef(mars_model$finalModel)
```

```
##           (Intercept)
##           -2.850834e+00
##           h(ap_hi-100)
##           9.199229e-02
##           h(16739-age_day)
##           -3.473246e-04
##           h(age_day-16739) * h(ap_hi-150)
##           -2.808598e-06
##           h(age_day-16739) * h(150-ap_hi)
##           6.204302e-06
##           cholesterolwell above normal
##           2.294506e+00
## h(ap_hi-100) * cholesterolwell above normal
##           -5.256094e-02
```

```
p1 = pdp::partial(mars_model, pred.var = c("age_day"), grid.resolution = 10) %>%
  autoplot()
p1
```

```
## Warning: Use of `object[[1L]]` is discouraged. Use `.data[[1L]]` instead.
```


```
## Warning: Use of `object[["yhat"]]` is discouraged. Use `.data[["yhat"]]`
## instead.
```



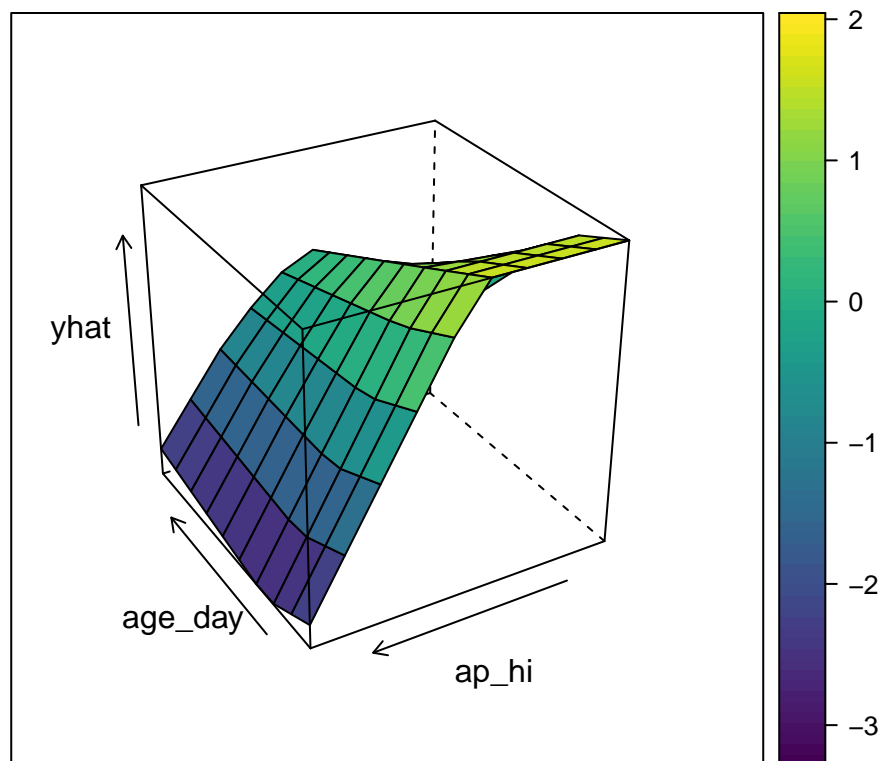
```
p2 = pdp::partial(mars_model, pred.var = c("ap_hi"), grid.resolution = 10) %>%
  autoplot()
p2
```

```
## Warning: Use of `object[[1L]]` is discouraged. Use `.data[[1L]]` instead.
```

```
## Warning: Use of `object[["yhat"]]` is discouraged. Use `.data[["yhat"]]`
## instead.
```



```
p3 = pdp::partial(mars_model, pred.var = c("age_day", "ap_hi"), grid.resolution = 10) %>%  
  pdp::plotPartial(levelplot = FALSE, zlab = "yhat", drape = TRUE, screen = list(z = 120, x = -60))  
p3
```





3. LDA

```
set.seed(2022)
lda_model = train(x = training_predictors_x,
                  y = training_outcome_y,
                  data = training_data,
                  method = "lda",
                  metric = "ROC",
                  trControl = control)
lda_model$results
```

```
## parameter      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1      none 0.7595429 0.7525714 0.6308571 0.05843652 0.08175705 0.0930469
```

```
lda_fit = lda(cardio ~., data = df_sample, subset = training_tag)
plot(lda_fit)
```

```
lda_fit$scaling
```

```
##                               LD1
## age_day                      0.0001587505
## height                     -0.0063718174
## weight                      0.0150851142
## ap_hi                       0.0429264353
## ap_lo                       0.0008343669
## gender1                     0.1539603773
## cholesterolabove normal    0.2990711209
## cholesterolwell above normal 0.8731256286
## glucabove normal          -0.0395037032
## glucwell above normal     0.1133729473
## smoke1                     -0.4623471382
## alco1                      -0.3063815769
## active1                    -0.1654022160
```

4. Boosting


```
set.seed(2022)
```

```
boost_grid = expand.grid(n.trees = c(1000, 2000, 3000, 4000),
                        interaction.depth = 1:6,
                        shrinkage = c(0.0004, 0.0006, 0.0008, 0.001),
                        n.minobsinnode = 1)
```

```
control_class = trainControl(method = "repeatedcv",
                             number = 10,
                             repeats = 5,
                             classProbs = TRUE,
                             summaryFunction = twoClassSummary
                             )

# Using caret perform boosting on the training data
boost_caret = train(cardio ~ .,
                    data = training_data,
                    method = "gbm",
                    tuneGrid = boost_grid,
                    trControl = control_class,
                    distribution = "adaboost",
                    metric = "ROC",
                    verbose = FALSE)


ggplot(boost_caret, highlight = TRUE)
```



```
boost_caret$bestTune
```

```
##      n.trees interaction.depth shrinkage n.minobsinnode
## 55      3000              2      8e-04              1
```

```
# Plot the variable importance
summary(boost_caret$finalModel, las = 2, cBars = 19, cex.names = 0.6)
```




```
##               var      rel.inf
## ap_hi         ap_hi 52.08437799
## age_day       age_day 21.50366721
## weight        weight 10.70296143
## ap_lo         ap_lo  7.11147104
## cholesterolwell above normal cholesterolwell above normal 5.31793794
## height        height  1.42389506
## active1       active1 0.85988309
## smoke1        smoke1 0.38443486
## glucabove normal glucabove normal 0.20070554
## alco1         alco1  0.13534147
## glucwell above normal glucwell above normal 0.13509642
## cholesterolabove normal cholesterolabove normal 0.10325443
## gender1       gender1 0.03697351
```

5. Classification Tree

```
set.seed(2022)
classification_tree_minMSE = rpart(formula = cardio ~ . ,
                                   data = training_data,
                                   control = rpart.control(cp = 0))
```

```
plotcp(classification_tree_minMSE)
```



```
# Obtain cp table
```

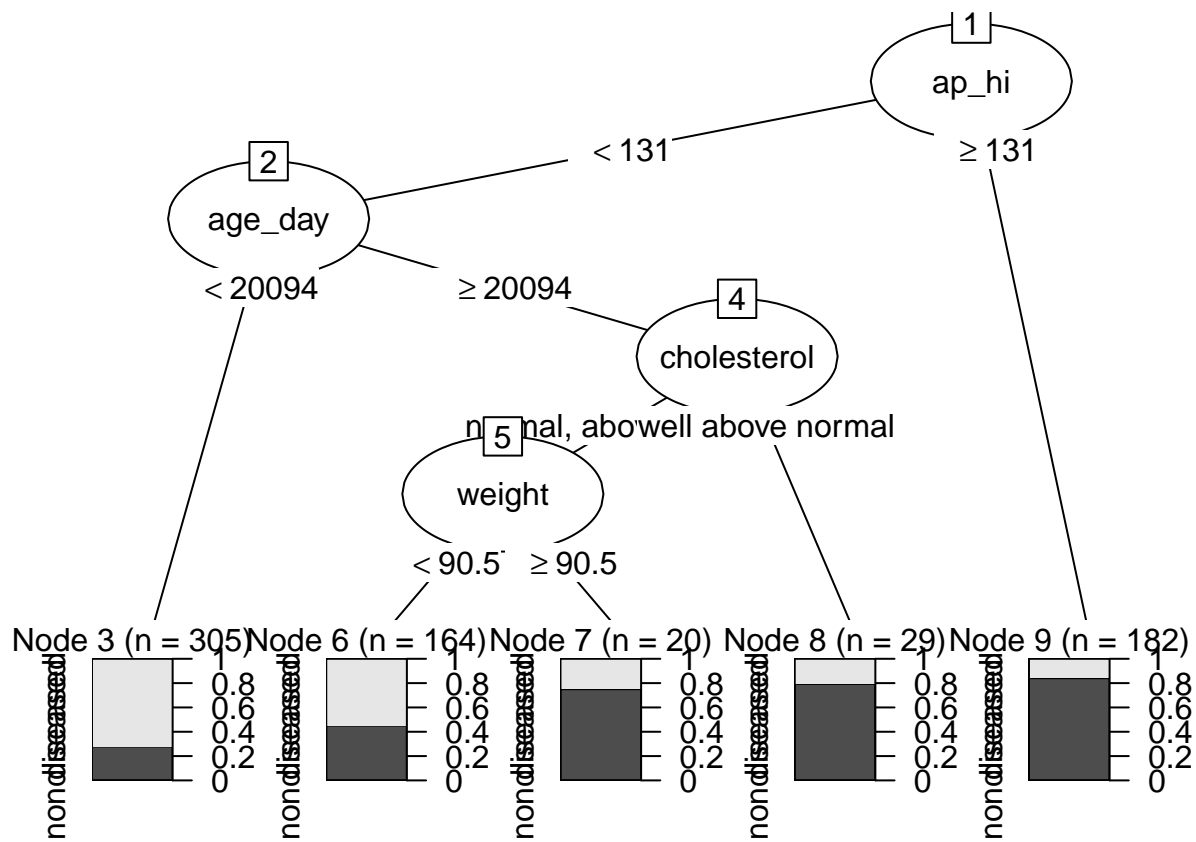
```
cp_table = printcp(classification_tree_minMSE)
```

```
##
## Classification tree:
## rpart(formula = cardio ~ ., data = training_data, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] active      age_day      alco         ap_hi        ap_lo        cholesterol
## [7] height      weight
##
## Root node error: 350/700 = 0.5
##
## n= 700
##
##      CP nsplit rel error  xerror    xstd
## 1  0.3542857     0  1.00000 1.10286 0.037596
## 2  0.0314286     1  0.64571 0.65143 0.035426
## 3  0.0228571     2  0.61429 0.68000 0.035809
## 4  0.0085714     4  0.56857 0.64571 0.035345
## 5  0.0071429     9  0.52000 0.66571 0.035622
## 6  0.0066667    11  0.50571 0.66286 0.035584
```

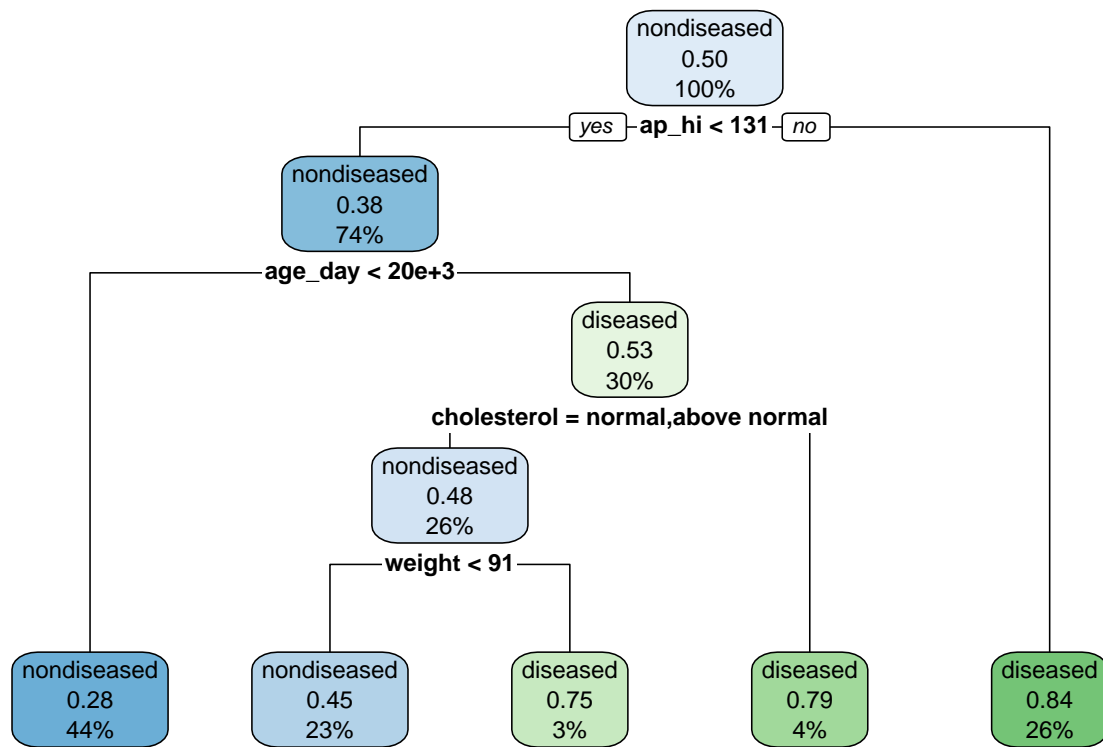
```
## 7  0.0057143    18  0.45143 0.66000 0.035545
## 8  0.0042857    19  0.44571 0.66000 0.035545
## 9  0.0028571    23  0.42857 0.65714 0.035506
## 10 0.0019048    24  0.42571 0.64857 0.035386
## 11 0.0000000    27  0.42000 0.68000 0.035809
```

```
df_MSE_min = which.min(cp_table[, 4])
final_class_tree_minMSE = prune(classification_tree_minMSE, cp = cp_table[df_MSE_min, 1])

# plot the minimum MSE classification tree
plot(as.party(final_class_tree_minMSE))
```



```
rpart.plot(final_class_tree_minMSE)
```



```

# Build classification tree using training dataset
classification_tree_1SE = prune(classification_tree_minMSE,
                                cp = cp_table[cp_table[, 4] < cp_table[df_MSE_min, 4] + cp_table[df_MSE_min, 4], 4])
plotcp(classification_tree_1SE)

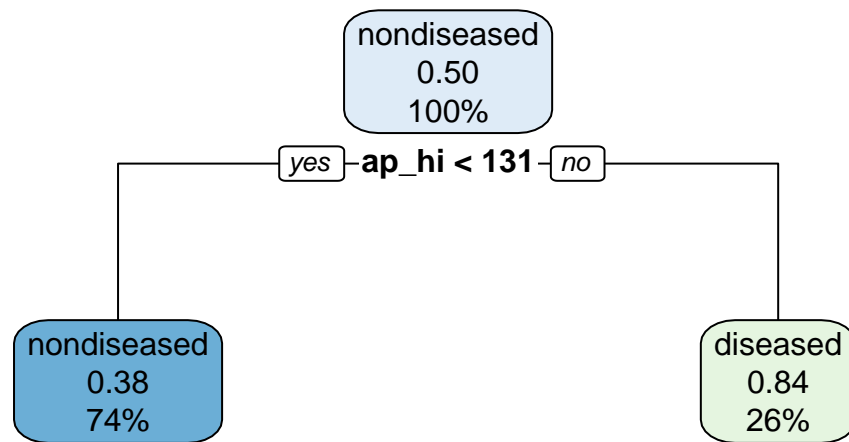
```




```
# Obtain cp table
cp_table_1se = printcp(classification_tree_1SE)

##
## Classification tree:
## rpart(formula = cardio ~ ., data = training_data, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] ap_hi
##
## Root node error: 350/700 = 0.5
##
## n= 700
##
##      CP nsplit rel error  xerror   xstd
## 1 0.354286      0  1.00000 1.10286 0.037596
## 2 0.031429      1  0.64571 0.65143 0.035426

# Plot the 1SE tree
rpart.plot(classification_tree_1SE)
```



```
plot(as.party(classification_tree_1SE))
```

```
auc(test_outcome_y, predict(classification_tree_1SE, newdata = test_data)[, 2])
```

```
## Setting levels: control = nondiseased, case = diseased
```

```
## Setting direction: controls < cases
```

```
## Area under the curve: 0.7133
```

```
ct_test_pred_prob = predict(classification_tree_1SE, newdata = test_data)[, 2]
```

```
ct_test_pred = rep("nondiseased", length(ct_test_pred_prob))
```

```
ct_test_pred[ct_test_pred_prob > 0.5] = "diseased"
```

```
confusionMatrix(data = as.factor(ct_test_pred),
                 reference = test_outcome_y
                 )
```

```
## Warning in confusionMatrix.default(data = as.factor(ct_test_pred), reference
## = test_outcome_y): Levels are not in the same order for reference and data.
## Refactoring data to match.
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##               Reference
## Prediction    nondiseased diseased
## nondiseased      135      71
## diseased         15      79
##
##               Accuracy : 0.7133
##               95% CI : (0.6586, 0.7638)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : 4.757e-14
##
##               Kappa : 0.4267
##
## Mcnemar's Test P-Value : 3.015e-09
##
##      Sensitivity : 0.9000
##      Specificity : 0.5267
##      Pos Pred Value : 0.6553
##      Neg Pred Value : 0.8404
##      Prevalence : 0.5000
##      Detection Rate : 0.4500
##      Detection Prevalence : 0.6867
##      Balanced Accuracy : 0.7133
##
##      'Positive' Class : nondiseased
##
```

Random forest

```
# Train caret random forest model
set.seed(2022)
# Grid of tuning parameters
rf_grid = expand.grid(mtry = 1:12,
                      splitrule = "gini",
                      min.node.size = seq(from = 2, to = 10, by = 2)
)

# Find best-fitting model after model fitting to optimize computational efficiency
rf_fit = train(cardio ~ .,
               data = training_data,
               method = "ranger",
               tuneGrid = rf_grid,
               metric = "ROC",
               trControl = control_class)

summary(rf_fit)
```

```
##               Length Class      Mode
## predictions      1400 -none-    numeric
## num.trees         1    -none-    numeric
## num.independent.variables 1 -none-    numeric
## mtry              1    -none-    numeric
## min.node.size     1    -none-    numeric
```

```
## prediction.error      1 -none-      numeric
## forest                10 ranger.forest list
## splitrule             1 -none-      character
## treetype              1 -none-      character
## call                  9 -none-      call
## importance.mode       1 -none-      character
## num.samples           1 -none-      numeric
## replace               1 -none-      logical
## xNames                13 -none-      character
## problemType           1 -none-      character
## tuneValue             3 data.frame list
## obsLevels             2 -none-      character
## param                 0 -none-      list
```

```
rf_pred = predict(rf_fit, newdata = test_data, type = "prob")[,2]
#rf_pred
```

```
#ConfusionMatrix
test_pred_rf = rep("nondiseased", length(rf_pred ))

test_pred_rf[rf_pred > 0.5] = "diseased"

confusionMatrix(data = as.factor(test_pred_rf),
                 reference = test_outcome_y
                 )
```

```
## Warning in confusionMatrix.default(data = as.factor(test_pred_rf), reference
## = test_outcome_y): Levels are not in the same order for reference and data.
## Refactoring data to match.
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  nondiseased diseased
##  nondiseased      116      37
##   diseased       34      113
##
##              Accuracy : 0.7633
##              95% CI : (0.7111, 0.8103)
##   No Information Rate : 0.5
##   P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.5267
##
##  Mcnemar's Test P-Value : 0.8124
##
##              Sensitivity : 0.7733
##              Specificity : 0.7533
##   Pos Pred Value : 0.7582
##   Neg Pred Value : 0.7687
##       Prevalence : 0.5000
##   Detection Rate : 0.3867
```

```
## Detection Prevalence : 0.5100
## Balanced Accuracy : 0.7633
##
## 'Positive' Class : nondiseased
##
```

```
auc(test_outcome_y, rf_pred)
```


```
## Setting levels: control = nondiseased, case = diseased
```

```
## Setting direction: controls < cases
```

```
## Area under the curve: 0.799
```


```
# Using impurity method to obtain variable importance
set.seed(2022)
rf_impurity_variable_importance = ranger(cardio ~ . ,
                                         data = training_data,
                                         mtry = rf_fit$bestTune[[1]],
                                         splitrule = "gini",
                                         min.node.size = rf_fit$bestTune[[3]],
                                         importance = "impurity")

# plot of variable importance using impurity
barplot(sort(ranger::importance(rf_impurity_variable_importance),
           decreasing = FALSE),
        las = 2,
        horiz = TRUE,
        cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan", "blue"))(19)
        )
```



```
# Using permutation method to obtain variable importance
rf_permutation_variable_importance = ranger(cardio ~ . ,
      data = training_data,
      mtry = rf_fit$bestTune[[1]],
      splitrule = "gini",
      min.node.size = rf_fit$bestTune[[3]],
      importance = "permutation",
      scale.permutation.importance = TRUE)

# plot of variable importance using permutation
barplot(sort(ranger::importance(rf_permutation_variable_importance),
      decreasing = FALSE),
      las = 2,
      horiz = TRUE,
      cex.names = 0.7,
      col = colorRampPalette(colors = c("cyan", "blue"))(19)
    )
```

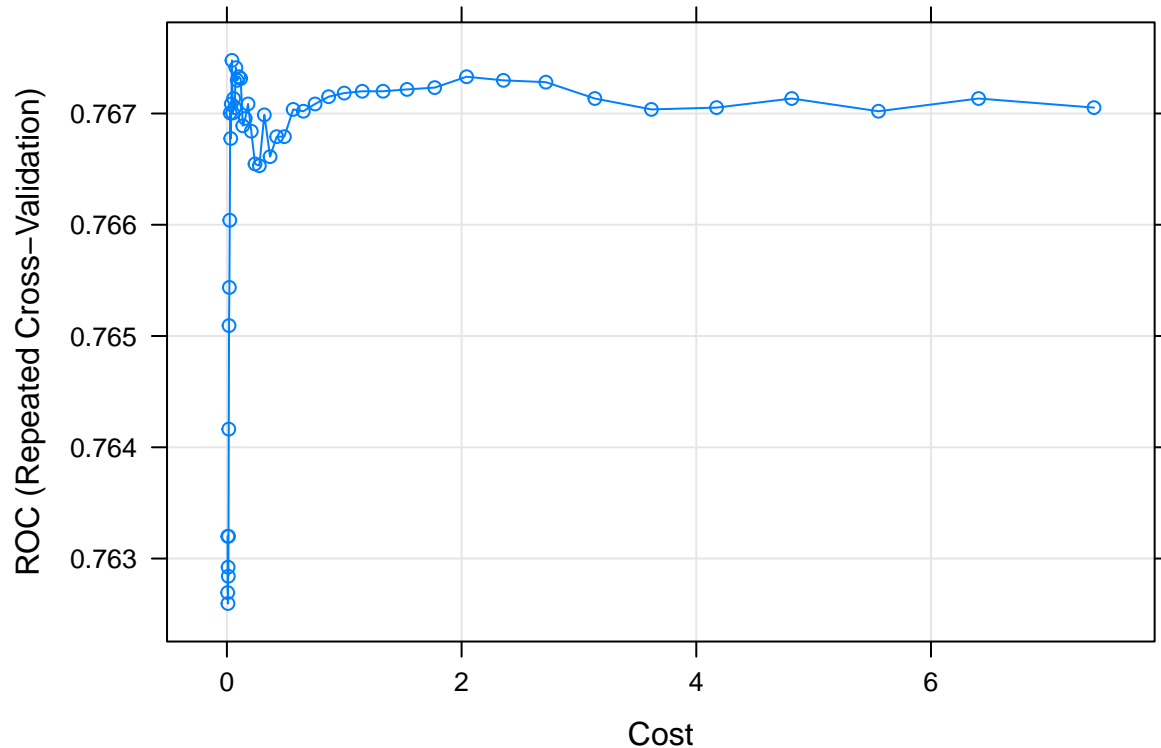


6. SVM

```
set.seed(2022)
linear_svc = train(cardio ~ .,
                    data = training_data,
                    method = "svmLinear",
                    tuneGrid = data.frame(C = exp(seq(-5, 2, len = 50))),
                    trControl = control_class,
                    scale = TRUE)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
plot(linear_svc)
```



```
linear_svc$bestTune
```

```
##           C
## 14 0.04315931
```

```
svm_pred = predict(linear_svc, newdata = test_data, type = "prob")[, 2]
test_pred_svm = rep("nondiseased", length(svm_pred))
test_pred_svm[svm_pred > 0.5] = "diseased"
confusionMatrix(data = as.factor(test_pred_svm),
                 reference = test_outcome_y)
```

```
## Warning in confusionMatrix.default(data = as.factor(test_pred_svm), reference
## = test_outcome_y): Levels are not in the same order for reference and data.
## Refactoring data to match.
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  nondiseased diseased
##  nondiseased      116      36
##   diseased       34      114
```

```
##
##          Accuracy : 0.7667
##          95% CI   : (0.7146, 0.8134)
##    No Information Rate : 0.5
##    P-Value [Acc > NIR] : <2e-16
##
##          Kappa : 0.5333
##
##    McNemar's Test P-Value : 0.9049
##
##          Sensitivity : 0.7733
##          Specificity : 0.7600
##    Pos Pred Value : 0.7632
##    Neg Pred Value : 0.7703
##          Prevalence : 0.5000
##    Detection Rate : 0.3867
##    Detection Prevalence : 0.5067
##    Balanced Accuracy : 0.7667
##
##    'Positive' Class : nondiseased
##
```

Final Model Selection:

```
set.seed(2022)

resamp = resamples(list(MARS = mars_model,
                        LDA = lda_model,
                        LOGISTIC = logistic_caret,
                        BOOSTING = boost_caret,
                        RANDOM_FOREST = rf_fit
                        #SVM = linear_svc
                        ))

summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: MARS, LDA, LOGISTIC, BOOSTING, RANDOM_FOREST
## Number of resamples: 50
##
## ROC
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## MARS      0.6106122 0.7358163 0.7714286 0.7701388 0.8077551 0.9134694    0
## LDA       0.6514286 0.7289796 0.7608163 0.7595429 0.8022449 0.8783673    0
## LOGISTIC  0.6579592 0.7383673 0.7669388 0.7655347 0.8018367 0.8832653    0
## BOOSTING  0.6432653 0.7442857 0.7775510 0.7810612 0.8189796 0.8946939    0
## RANDOM_FOREST 0.6865306 0.7314286 0.7881633 0.7795755 0.8126531 0.9093878    0
##
## Sens
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## MARS      0.5714286 0.6571429 0.7428571 0.7302857 0.7928571 0.9142857    0
```



```
## LDA          0.6285714 0.6857143 0.7428571 0.7525714 0.8000000 0.9428571 0
## LOGISTIC     0.6285714 0.6928571 0.7428571 0.7571429 0.8000000 0.9428571 0
## BOOSTING     0.6285714 0.7214286 0.7714286 0.7714286 0.8285714 0.9428571 0
## RANDOM_FOREST 0.5428571 0.6928571 0.7714286 0.7582857 0.8000000 0.9428571 0
##
## Spec
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## MARS       0.4571429 0.6000000 0.6571429 0.6640000 0.7142857 0.8285714 0
## LDA        0.4000000 0.5714286 0.6285714 0.6308571 0.6857143 0.8285714 0
## LOGISTIC   0.3714286 0.6000000 0.6428571 0.6491429 0.7142857 0.8285714 0
## BOOSTING   0.4285714 0.6000000 0.6285714 0.6405714 0.7071429 0.8571429 0
## RANDOM_FOREST 0.4285714 0.6000000 0.6714286 0.6594286 0.7142857 0.8000000 0
```

```
bwplot(resamp)
```

